

Linear Regression

→ feature - x , label - $y \rightarrow h_\theta(x) \rightarrow$ to be predicted as y .

→ $h_\theta(x) = \theta_0 + \theta_1 x$, where θ are parameters (hypothesis)

→ Cost function: $\frac{1}{2m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2 = J(\theta_0, \theta_1)$

This is called the squared error function.

→ Gradient Descent:

repeat until convergence

$$\left\{ \begin{array}{l} \theta_i := \theta_i - \alpha \frac{\partial}{\partial \theta_i} J(\theta_0, \dots, \theta_1) \quad [\text{Simultaneously for all } i] \end{array} \right\}$$

→ here,

repeat until convergence

$$\left\{ \begin{array}{l} \theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \\ \theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \cdot x^{(i)} \end{array} \right\} \quad \text{Simultaneously}$$

< This is written after calculating derivatives manually >

→ The cost function is convex, so it does not matter where we start. So we can randomly initialize. It will converge to same point.

↳ Randomly initialize → train using gradient descent (batch)

GET MODEL